

Track experiments

[Try in a Colab Notebook here →](#)

Rapid experimentation is fundamental to machine learning. [In this tutorial, we use W&B to track and visualize experiments so that we can quickly iterate and understand our results.](#)

A shared dashboard for your experiments

With just a few lines of code, you'll get rich, interactive, shareable dashboards [which you can see yourself here](#).

Data & Privacy

We take security very seriously, and our cloud-hosted dashboard uses industry standard best practices for encryption. If you're working with datasets that cannot leave your enterprise cluster, we have [on-prem](#) installations available.

It's also easy to download all your data and export it to other tools — like custom analysis in a Jupyter notebook. Here's [more on our API](#).

□ Install `wandb` library and login

Start by installing the library and logging in to your free account.

```
!pip install wandb -qU
```

```
# Log in to your W&B account
import wandb
wandb.login()
```

Run an experiment

1. [Start a new run and pass in hyperparameters to track](#)
2. [Log metrics from training or evaluation](#)
3. [Visualize results in the dashboard](#)

```

import random

# Launch 5 simulated experiments
total_runs = 5
for run in range(total_runs):
    # 🐛 1 Start a new run to track this script
    wandb.init(
        # Set the project where this run will be logged
        project="basic-intro",
        # We pass a run name (otherwise it'll be randomly assigned, like sunshine-
        lollypop-10)
        name=f"experiment_{run}",
        # Track hyperparameters and run metadata
        config={
            "learning_rate": 0.02,
            "architecture": "CNN",
            "dataset": "CIFAR-100",
            "epochs": 10,
        })

    # This simple block simulates a training loop logging metrics
    epochs = 10
    offset = random.random() / 5
    for epoch in range(2, epochs):
        acc = 1 - 2 ** -epoch - random.random() / epoch - offset
        loss = 2 ** -epoch + random.random() / epoch + offset

        # 🐛 2 Log metrics from your script to W&B
        wandb.log({"acc": acc, "loss": loss})

    # Mark the run as finished
    wandb.finish()

```

3 When you run this code, you can find your interactive dashboard by clicking any of the 🖱️ wandb links above.

🔥 Simple Pytorch Neural Network

👉 Run this model to train a simple MNIST classifier, and click on the project page link to see your results stream in live to a W&B project.

Any run in `wandb` automatically logs [metrics](#), [system information](#), [hyperparameters](#), [terminal output](#) and you'll see an [interactive table](#) with model inputs and outputs.

Set up Dataloader

To run this example, we'll need to install PyTorch. If you're using Google Colab, it is already preinstalled.

```
!pip install torch torchvision
```

```
import wandb
import math
import random
import torch, torchvision
import torch.nn as nn
import torchvision.transforms as T

device = "cuda:0" if torch.cuda.is_available() else "cpu"

def get_dataloader(is_train, batch_size, slice=5):
    "Get a training dataloader"
    full_dataset = torchvision.datasets.MNIST(root=".", train=is_train,
transform=T.ToTensor(), download=True)
    sub_dataset = torch.utils.data.Subset(full_dataset, indices=range(0,
len(full_dataset), slice))
    loader = torch.utils.data.DataLoader(dataset=sub_dataset,
batch_size=batch_size,
shuffle=True if is_train else False,
pin_memory=True, num_workers=2)

    return loader

def get_model(dropout):
    "A simple model"
    model = nn.Sequential(nn.Flatten(),
nn.Linear(28*28, 256),
nn.BatchNorm1d(256),
nn.ReLU(),
nn.Dropout(dropout),
nn.Linear(256,10)).to(device)

    return model

def validate_model(model, valid_dl, loss_func, log_images=False, batch_idx=0):
    "Compute performance of the model on the validation dataset and log a wandb.Table"
    model.eval()
    val_loss = 0.
    with torch.inference_mode():
        correct = 0
        for i, (images, labels) in enumerate(valid_dl):
            images, labels = images.to(device), labels.to(device)

            # Forward pass →
            outputs = model(images)
            val_loss += loss_func(outputs, labels)*labels.size(0)

            # Compute accuracy and accumulate
            _, predicted = torch.max(outputs.data, 1)
            correct += (predicted == labels).sum().item()

        # Log one batch of images to the dashboard, always same batch_idx.
```

```

        if i==batch_idx and log_images:
            log_image_table(images, predicted, labels, outputs.softmax(dim=1))
    return val_loss / len(valid_dl.dataset), correct / len(valid_dl.dataset)

```

```

def log_image_table(images, predicted, labels, probs):
    "Log a wandb.Table with (img, pred, target, scores)"
    # 🐛 Create a wandb Table to Log images, labels and predictions to
    table = wandb.Table(columns=["image", "pred", "target"]+[f"score_{i}" for i in
range(10)])
    for img, pred, targ, prob in zip(images.to("cpu"), predicted.to("cpu"),
labels.to("cpu"), probs.to("cpu")):
        table.add_data(wandb.Image(img[0].numpy()*255), pred, targ, *prob.numpy())
    wandb.log({"predictions_table":table}, commit=False)

```

Train Your Model

Launch 5 experiments, trying different dropout rates

```
for _ in range(5):
```

🐛 initialise a wandb run

```

wandb.init(
    project="pytorch-intro",
    config={
        "epochs": 10,
        "batch_size": 128,
        "lr": 1e-3,
        "dropout": random.uniform(0.01, 0.80),
    })

```

Copy your config

```
config = wandb.config
```

Get the data

```

train_dl = get_dataloader(is_train=True, batch_size=config.batch_size)
valid_dl = get_dataloader(is_train=False, batch_size=2*config.batch_size)
n_steps_per_epoch = math.ceil(len(train_dl.dataset) / config.batch_size)

```

A simple MLP model

```
model = get_model(config.dropout)
```

Make the Loss and optimizer

```

loss_func = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=config.lr)

```

Training

```
example_ct = 0
```

```
step_ct = 0
```

```
for epoch in range(config.epochs):
```

```
    model.train()
```

```
    for step, (images, labels) in enumerate(train_dl):
```

```
        images, labels = images.to(device), labels.to(device)
```

```

        outputs = model(images)
        train_loss = loss_func(outputs, labels)
        optimizer.zero_grad()
        train_loss.backward()
        optimizer.step()

        example_ct += len(images)
        metrics = {"train/train_loss": train_loss,
                  "train/epoch": (step + 1 + (n_steps_per_epoch * epoch)) /
n_steps_per_epoch,
                  "train/example_ct": example_ct}

        if step + 1 < n_steps_per_epoch:
            # 🐛 Log train metrics to wandb
            wandb.log(metrics)

        step_ct += 1

    val_loss, accuracy = validate_model(model, valid_dl, loss_func, log_images=
(epoch==(config.epochs-1)))

    # 🐛 Log train and validation metrics to wandb
    val_metrics = {"val/val_loss": val_loss,
                  "val/val_accuracy": accuracy}
    wandb.log(**metrics, **val_metrics)

    print(f"Train Loss: {train_loss:.3f}, Valid Loss: {val_loss:.3f}, Accuracy:
{accuracy:.2f}")

    # If you had a test set, this is how you could log it as a Summary metric
    wandb.summary['test_accuracy'] = 0.8

    # 🐛 Close your wandb run
    wandb.finish()

```

You have now trained your first model using wandb! 🖱️ Click on the wandb link above to see your metrics

🔔 Try W&B Alerts

W&B Alerts allows you to send alerts, triggered from your Python code, to your Slack or email. There are 2 steps to follow the first time you'd like to send a Slack or email alert, triggered from your code:

- 1) [Turn on Alerts in your W&B User Settings](#)
- 2) Add `wandb.alert()` to your code:

```

wandb.alert(
    title="Low accuracy",

```

```
text=f"Accuracy is below the acceptable threshold"
)
```

See the minimal example below to see how to use `wandb.alert`. You can find the full docs for [W&B Alerts here](#)

```
# Start a wandb run
wandb.init(project="pytorch-intro")
```

```
# Simulating a model training loop
acc_threshold = 0.3
for training_step in range(1000):
```

```
    # Generate a random number for accuracy
    accuracy = round(random.random() + random.random(), 3)
    print(f'Accuracy is: {accuracy}, {acc_threshold}')
```

```
    # 🐛 Log accuracy to wandb
    wandb.log({"Accuracy": accuracy})
```

```
    # 🛑 If the accuracy is below the threshold, fire a W&B Alert and stop the run
    if accuracy <= acc_threshold:
        # 🐛 Send the wandb Alert
        wandb.alert(
            title='Low Accuracy',
            text=f'Accuracy {accuracy} at step {training_step} is below the acceptable
threshold, {acc_threshold}',
        )
        print('Alert triggered')
        break
```

```
    # Mark the run as finished (useful in Jupyter notebooks)
    wandb.finish()
```

What's next?

The next tutorial, you will learn how to view & analyze model predictions using W&B Tables:

👉 View & Analyze Model Predictions

Was this page helpful? 👍 👎